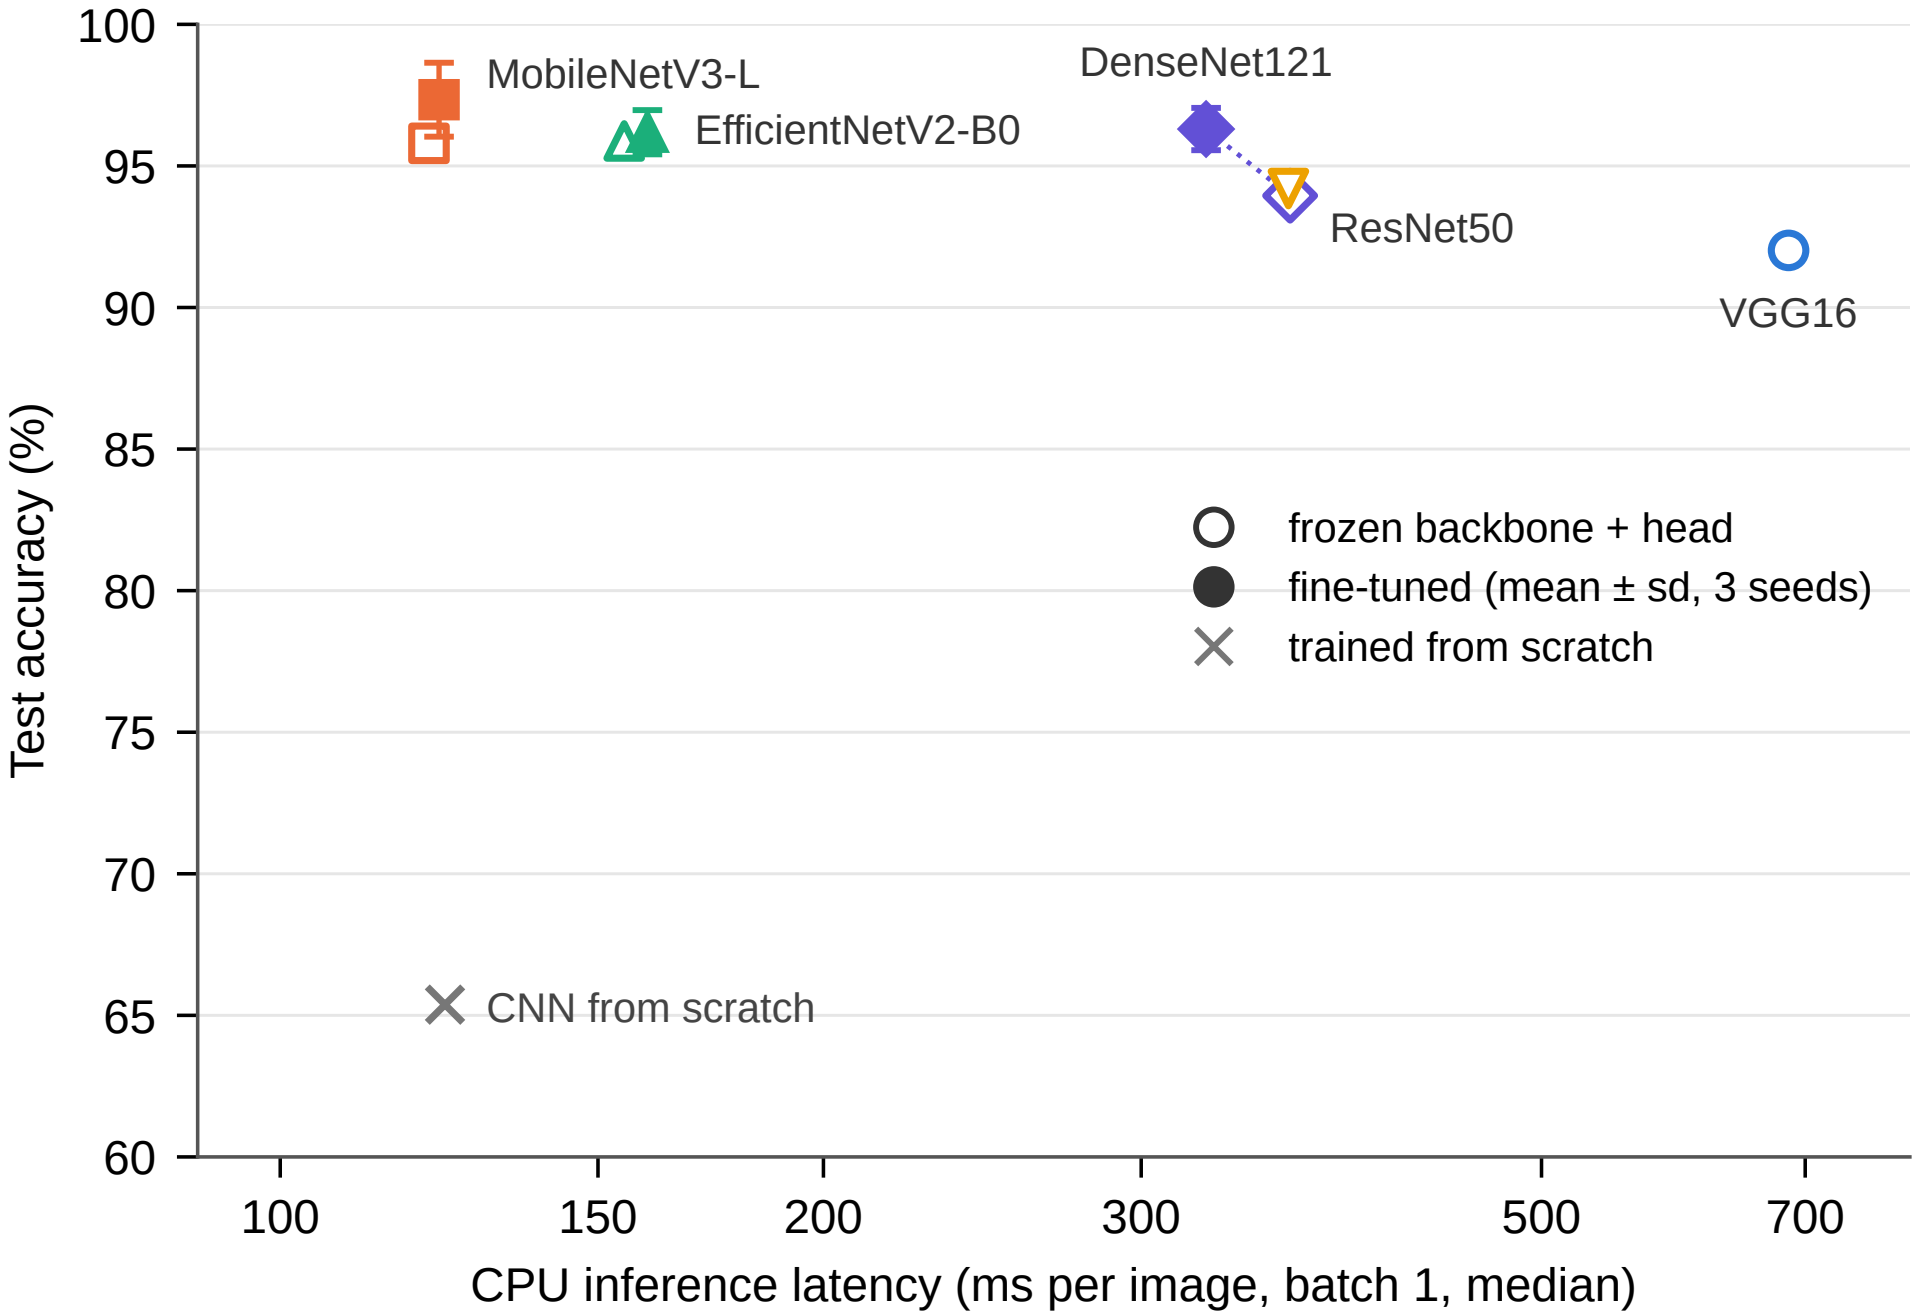
100
95
90
85
80
75
70
65
60
Test accuracy (%)
MobileNetV3-L
EfficientNetV2-B0
DenseNet121
ResNet50
VGG16
CNN from scratch
frozen backbone + head
fine-tuned (mean ± sd, 3 seeds)
trained from scratch
100
150
200
300
500
700
CPU inference latency (ms per image, batch 1, median)